

Final Project

Hannah Wilder and Chathura Gunasekara

April 9, 2016

Notes: Possible source of population data: <http://www.abs.gov.au/AUSSTATS/abs@.nsf/DetailsPage/3105.0.65.0012014?OpenDocument>

Change working directory here

Load data (assumes file is in working directory)


```
#load the data
beerData<-read.csv("monthly-beer-production-in-austr.csv")

#cut off the last row which is NA
beerData<-beerData[~nrow(beerData),]
colnames(beerData)<-c("Month", "Production")


#turn into time series also hold back the last two years of data for forecasting
beerTS<-ts(beerData[1:(nrow(beerData)-24),2], frequency=12, start=c(1956,1))
beer_forecast<-ts(beerData[(nrow(beerData)-23):nrow(beerData), 2], start=c(1993,9), frequency=12)
```

Plot data

```
par(mfrow=c(1,1))
plot(beerTS, main="Beer Production in Australia by Month")
```




```
plot(beerTS, main="Beer Production in Australia by Month (seasons marked)", type="l")
points(y=beerTS, x=time(beerTS), pch=as.vector(season(beerTS)))
```



Another plot to show seasonality

```
require(fpp)
seasonplot(beerTS,year.labels=TRUE,ylab="megaliters",main="Seasonal plot: quarterly beer production", col=rain)
```

Seasonal plot: quarterly beer production




In the plot we see obvious seasonality with higher production in November and December and lower production in June and July. There is a trend which may be difficult to fit as it doesn't appear to be a "well known" function like a linear or quadratic function, so we'll have to experiment. It also looks like the variance of the data is larger in the middle, so we will probably want to take the log of our data to correct that variance issue.

Decomposing the time series to see trends and patterns

```
decompbeer = decompose (beerTS, type="additive")  
plot (decompbeer)
```

Decomposition of additive time series



by looking at the decomposed figures, i was wondering what if we plot a harmonic function with a quadratic polynomial... like imposing a sine curve with 2nd order poly ?

Investigate possible relationship with population data

```
#load population data  
library(reshape)
```

```
## Warning: package 'reshape' was built under R version 3.2.5
```

```
library(ggplot2)
```

```
## Warning: package 'ggplot2' was built under R version 3.2.5
```

```
#Clean up population data  
pop_totalData<-t(read.csv("Pop_total.csv", row.names=1))  
dropCols<-colnames(pop_totalData) %in% c("Unspecified","Period not indicated")  
rownames(pop_totalData)<-c(1921:2011)  
pop_totalDataLong<-pop_totalData[,!dropCols]  
pop_totalData<-pop_totalData[paste(1956:1995),!dropCols]  
  
#Aggregate beer data  
beerYear<-seq(from=1956, to=1996, by=1)  
beerYear<-rep(beerYear, each=12)  
beerYear<-beerYear[1:nrow(beerData)]  
beerAg<-aggregate(beerData[,2], FUN=mean, by=list(year=beerYear))  
  
#Attach to beer data
```

```

beerPop<-data.frame(cbind(beer=beerAg[,2],pop_totalData))
beerPopScale<-scale(beerPop)


beerPopRes<-melt(beerPopScale, variable.name="series")
colnames(beerPopRes)<-c("time", "series", "stddev")

allNames<-colnames(beerPop)[2:length(colnames(beerPop))]


#Plot data for each age group and beer data on same plot
par(mfrow=c(2,2))
for (name in allNames) {
  subset_data<-subset(beerPopRes, beerPopRes$series%in%c("beer", name))
  newPlot<-ggplot(subset_data, aes(time,stddev)) + geom_line(aes(colour = series)) +ggtitle(paste("Pop (age gr", name, ") and Beer Prod"))
  print(newPlot)
}

```


Pop (age group X0.4) and Beer Prod




Pop (age group X5.9) and Beer Prod




Pop (age group X10.14) and Beer Prod




Pop (age group X15.19) and Beer Prod




Pop (age group X20.24) and Beer Prod




Pop (age group X25.29) and Beer Prod




Pop (age group X30.34) and Beer Prod




Pop (age group X35.39) and Beer Prod




Pop (age group X40.44) and Beer Prod




Pop (age group X45.49) and Beer Prod




Pop (age group X50.54) and Beer Prod




Pop (age group X55.59) and Beer Prod




Pop (age group X60.64) and Beer Prod




Pop (age group X65.69) and Beer Prod




Pop (age group X70.74) and Beer Prod




Pop (age group X75.79) and Beer Prod




Pop (age group X80.84) and Beer Prod



(age group X85.and.over) and Beer Prod



Pop (age group Total) and Beer Prod



```
par(mfrow=c(1,1))
```


```
#Make a model based on the 15-19 age group
yearModel1<-lm(beer ~ X15.19, data=beerPop)
```

It appears that there may be a relationship with the 15-19 age block, which makes sense since the legal drinking age is 18. We may be able to use this to remove some of our trend. However, the 10-14 age group numbers look like they might have potential if shifted forward a few years. This makes sense since these children will grow up and start drinking beer.

```
#Explore lagged x10.14 data
laggedData<-data.frame(beer=beerAg[,2])
models<-list()
modelRsqr<-c()
for (lag in 0:8) {
  newColNames<-c(colnames(laggedData), paste("lag", lag, sep=""))
  newLag<-pop_totalDataLong[paste(1956:1995-lag, "10-14")]
  laggedData<-data.frame(laggedData, newLag)
  newModel<-lm(beer ~ newLag, data=laggedData)
  models[[paste("lag", lag, sep="")]]<-newModel
  modelRsqr<-c(modelRsqr, summary(newModel)$r.squared)
  colnames(laggedData)<-newColNames
}
```


```
}
```

```
plot(modelRsqr, main="R Squared Values for lags of 10-14 age group", xlab="lag number", ylab="R squared value")  
abline(v=6, lty=2)
```




```
lagDataScale<-scale(laggedData)[,c(1,8)]  
lagDataMelt<-melt(lagDataScale, variable.name="series")  
colnames(lagDataMelt)<-c("time", "series", "stddev")  
  
newPlot<-ggplot(lagDataMelt, aes(time,stddev)) + geom_line(aes(colour = series)) +ggtitle(paste("Lags of 10-14  
print(newPlot)
```


Lags of 10–14 Pop and Beer Prod



```
plot(residuals(models[["lag6"]]), main="Residuals from modeling beer production with 10-14 lag 6")
```

Residuals from modeling beer production with 10–14 lag 6



We see that lag 6 is the optimal lag in terms of R-squared values. This makes sense because in 6 years, this age group will be 16-20, or right around drinking age. We can see in the residuals that it isn't perfect, but this pattern may be easier to model

than what we had before, it looks much more like a regular polynomial.

Try to figure out deterministic trend


```
t<-1:length(beerTS)
t2<-t^2
t3<-t^3
t4<-t^4
t5<-t^5

quadFit<-lm(beerTS~t+t2)
summary(quadFit)
```

```
##
## Call:
## lm(formula = beerTS ~ t + t2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -46.861 -14.133  -1.991   11.937   61.174
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  6.560e+01  2.828e+00   23.20  <2e-16 ***
## t            5.429e-01  2.883e-02   18.83  <2e-16 ***
## t2          -7.721e-04  6.163e-05  -12.53  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.95 on 449 degrees of freedom
## Multiple R-squared:  0.6616, Adjusted R-squared:  0.6601
## F-statistic:  439 on 2 and 449 DF,  p-value: < 2.2e-16
```


```
#### plot the data and the fitted quadratic trend function
plot(x=1:length(beerTS),y=beerTS,type='o',ylab="",xlab="Time - Number of Months Since Jan 1956",main="Quadratic trend",
curve(expr = coef(quadFit)[1]+coef(quadFit)[2]*x+coef(quadFit)[3]*x^2+coef(quadFit)[4]*x^3,lty=1,add = TRUE, col="red",lty=1))
```

Quadratic Fit on Beer Production Data



```
quadFit_resid<-ts(residuals(quadFit),frequency=12, start=c(1956,1))  
plot(quadFit_resid, main="Residuals from a Quadratic Trend Fit")  
abline(h=0)
```

Residuals from a Quadratic Trend Fit



```

cubicFit<-lm(beerTS~t+t2+t3)
summary(cubicFit)

##
## Call:
## lm(formula = beerTS ~ t + t2 + t3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -50.660 -13.783  -2.601  12.434  57.639
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  7.745e+01  3.695e+00  20.963  < 2e-16 ***
## t            2.307e-01  7.056e-02   3.270  0.00116 **
## t2           9.490e-04  3.617e-04   2.624  0.00900 **
## t3          -2.533e-06  5.249e-07  -4.826  1.92e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.47 on 448 degrees of freedom
## Multiple R-squared:  0.6784, Adjusted R-squared:  0.6762
## F-statistic: 315 on 3 and 448 DF, p-value: < 2.2e-16


```

```

#### plot the data and the fitted quadratic trend function
plot(x=1:length(beerTS),y=beerTS,type='o',ylab="",xlab="Time - Number of Months Since Jan 1956",main="Cubic Fit on Beer Production Data",add = TRUE)
curve(expr = coef(cubicFit)[1]+coef(cubicFit)[2]*x+coef(cubicFit)[3]*x^2+coef(cubicFit)[4]*x^3,lty=1,add = TRUE)

```

Cubic Fit on Beer Production Data




```

cubicFit_resid<-ts(residuals(cubicFit),frequency=12, start=c(1956,1))
plot(cubicFit_resid, main="Residuals from a Cubic Trend Fit")
abline(h=0)

```

Residuals from a Cubic Trend Fit




```
order4polyFit<-lm(beerTS~t+t2+t3+t4)
summary(order4polyFit)
```

```
##
## Call:
## lm(formula = beerTS ~ t + t2 + t3 + t4)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -50.079 -12.721  -3.199  10.135  57.983
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  9.037e+01  4.536e+00  19.924  < 2e-16 ***
## t           -3.341e-01  1.384e-01  -2.414   0.0162 *
## t2            6.545e-03  1.241e-03   5.276  2.07e-07 ***
## t3           -2.173e-05  4.113e-06  -5.285  1.97e-07 ***
## t4            2.119e-08  4.504e-09   4.706  3.38e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.03 on 447 degrees of freedom
## Multiple R-squared:  0.6935, Adjusted R-squared:  0.6908
## F-statistic: 252.9 on 4 and 447 DF, p-value: < 2.2e-16
```

```
#### plot the data and the fitted 4th order polynomial trend function
```


```
plot(x=1:length(beerTS),y=beerTS,type='o',ylab="",xlab="Time - Number of Months Since Jan 1956",main="order4po
curve(expr = coef(order4polyFit)[1]+coef(order4polyFit)[2]*x+coef(order4polyFit)[3]*x^2+coef(order4polyFit)[4]
```

order4poly Fit on Beer Production Data



```
order4polyFit_resid<-ts(residuals(order4polyFit),frequency=12, start=c(1956,1))  
plot(order4polyFit_resid, main="Residuals from a order4poly Trend Fit")  
abline(h=0)
```

Residuals from a order4poly Trend Fit



```
order5polyFit<-lm(beerTS~t+t2+t3+t4+t5)
summary(order5polyFit)
```

```
##
## Call:
## lm(formula = beerTS ~ t + t2 + t3 + t4 + t5)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-50.069	-12.729	-3.179	10.132	58.012

```
##
## Coefficients:
```


	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	9.029e+01	5.483e+00	16.469	<2e-16 ***
t	-3.288e-01	2.436e-01	-1.350	0.1778
t2	6.463e-03	3.323e-03	1.945	0.0524 .
t3	-2.126e-05	1.858e-05	-1.144	0.2532
t4	2.000e-08	4.519e-08	0.443	0.6582
t5	1.049e-12	3.970e-11	0.026	0.9789

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.05 on 446 degrees of freedom
## Multiple R-squared:  0.6935, Adjusted R-squared:  0.6901
## F-statistic: 201.9 on 5 and 446 DF,  p-value: < 2.2e-16
```


```
#### plot the data and the fitted 5th order polynomial trend function
```

```
plot(x=1:length(beerTS),y=beerTS,type='o',ylab="",xlab="Time - Number of Months Since Jan 1956",main="order5poly Fit on Beer Production Data")
curve(expr = coef(order5polyFit)[1]+coef(order5polyFit)[2]*x+coef(order5polyFit)[3]*x^2+coef(order5polyFit)[4]*x^3+coef(order5polyFit)[5]*x^4+coef(order5polyFit)[6]*x^5,lty=2)
```

order5poly Fit on Beer Production Data



```
order5polyFit_resid<-ts(residuals(order5polyFit),frequency=12, start=c(1956,1))
plot(order5polyFit_resid, main="Residuals from a order5poly Trend Fit")
abline(h=0)
```



It looks like a 4th order polynomial might take care of the worst of it, the question is are we okay with using a 4th order polynomial or should we drop it down to a cubic function and just deal with it? I found population data and I would be interested to see if we can find a good correlation there (total population won't work, I already looked at that, but maybe a specific age group?)

Assume we go with the 4th order polynomial for now. Let's see what we can do about the seasonality with a seasonal means model

```
library(TSA)
month=season(order4polyFit_resid)
seasMeansModel<-lm(order4polyFit_resid~month)
summary(seasMeansModel)
```


```
##
## Call:
## lm(formula = order4polyFit_resid ~ month)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -37.789  -6.263   0.327   6.128  35.453
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      4.108      1.652   2.487 0.013268 *
## monthFebruary    -8.810      2.336  -3.771 0.000185 ***
## monthMarch        1.929      2.336   0.826 0.409372
## monthApril       -11.085      2.336  -4.744 2.83e-06 ***
## monthMay         -15.540      2.336  -6.651 8.65e-11 ***
```



```
## monthJune      -29.051      2.336 -12.434 < 2e-16 ***
## monthJuly      -19.403      2.336  -8.304 1.25e-15 ***
## monthAugust    -13.661      2.336  -5.847 9.78e-09 ***
## monthSeptember -10.151      2.352  -4.316 1.97e-05 ***
## monthOctober     9.861       2.352   4.192 3.34e-05 ***
## monthNovember   17.904      2.352   7.612 1.66e-13 ***
## monthDecember   30.406      2.352  12.927 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 10.18 on 440 degrees of freedom
## Multiple R-squared:  0.7181, Adjusted R-squared:  0.7111
## F-statistic: 101.9 on 11 and 440 DF,  p-value: < 2.2e-16
```

```
seasMeansModel_resid<-ts(residuals(seasMeansModel),frequency=12, start=c(1956,1))
plot(seasMeansModel_resid, main="Residuals from Seasonal Means Model \n(after fitting 4th order polynomial)")
abline(h=0)
```

Residuals from Seasonal Means Model (after fitting 4th order polynomial)



With an adjusted R-squared value of 71%, this is looking pretty good, but in the residual plot you can still see the variance increasing over time. In addition, there is a noticeable “wave” in the residuals that starts around 1970, but I’m not sure what to do about that yet. For now, let’s go back, log the data, and apply both the 4th order polynomial and the seasonal means model at the same time.


```
logBeer<-log(beerTS)
t<-1:length(logBeer)
t2<-t^2
t3<-t^3
t4<-t^4
month<-season(logBeer)

logSeasPoly<-lm(logBeer~t+t2+t3+t4+month)
summary(logSeasPoly)
```

```
##
## Call:
## lm(formula = logBeer ~ t + t2 + t3 + t4 + month)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.262750 -0.039816  0.003297  0.043475  0.184483
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   4.506e+00  1.945e-02 231.675 < 2e-16 ***
## t             -1.796e-03  5.014e-04  -3.583 0.000378 ***
## t2             4.926e-05  4.494e-06   10.962 < 2e-16 ***
## t3            -1.745e-07  1.490e-08  -11.713 < 2e-16 ***
## t4             1.785e-10  1.631e-11   10.941 < 2e-16 ***
## monthFebruary -6.602e-02  1.580e-02  -4.178 3.56e-05 ***
## monthMarch     1.069e-02  1.580e-02    0.676 0.499268
## monthApril    -8.834e-02  1.581e-02  -5.590 4.02e-08 ***
## monthMay      -1.271e-01  1.581e-02  -8.042 8.39e-15 ***
## monthJune     -2.427e-01  1.581e-02 -15.354 < 2e-16 ***
## monthJuly     -1.578e-01  1.581e-02  -9.984 < 2e-16 ***
## monthAugust   -1.089e-01  1.581e-02  -6.890 1.96e-11 ***
## monthSeptember -7.261e-02  1.591e-02  -4.563 6.57e-06 ***
## monthOctober   6.706e-02  1.591e-02    4.213 3.06e-05 ***
## monthNovember  1.172e-01  1.592e-02    7.363 9.03e-13 ***
## monthDecember  1.936e-01  1.592e-02   12.164 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.06889 on 436 degrees of freedom
## Multiple R-squared:  0.9365, Adjusted R-squared:  0.9344
## F-statistic: 429 on 15 and 436 DF, p-value: < 2.2e-16
```

```
logSeasPoly_resid<-ts(residuals(logSeasPoly),frequency=12, start=c(1956,1))
plot(logSeasPoly_resid, main="Residuals from Logged Beer\nseasonal Means and 4th order poly fit at same time",
points(y=logSeasPoly_resid, x=time(logSeasPoly_resid), pch=as.vector(season(logSeasPoly_resid)))
abline(h=0)
```

Residuals from Logged Beer seasonal Means and 4th order poly fit at same time



Let's take a look and see if we have a stationary series yet

```
# d
adf.test(logSeasPoly_resid)

## Warning in adf.test(logSeasPoly_resid): p-value smaller than printed p-
## value

##
## Augmented Dickey-Fuller Test
##
## data: logSeasPoly_resid
## Dickey-Fuller = -5.3999, Lag order = 7, p-value = 0.01
## alternative hypothesis: stationary


pp.test(logSeasPoly_resid)

## Warning in pp.test(logSeasPoly_resid): p-value smaller than printed p-value

##
## Phillips-Perron Unit Root Test
##
## data: logSeasPoly_resid
## Dickey-Fuller Z(alpha) = -489.81, Truncation lag parameter = 5,
## p-value = 0.01
## alternative hypothesis: stationary


# p & q
par(mfrow=c(1,1))
acf(logSeasPoly_resid, lag.max=100)
```

Series logSeasPoly_resid



```
pacf(logSeasPoly_resid, lag.max=100)
```

Series logSeasPoly_resid



```
par(mfrow=c(1,1))
eacf(logSeasPoly_resid)
```

```
## AR/MA
##   0 1 2 3 4 5 6 7 8 9 10 11 12 13
## 0 o o x o x x x x x o x x o x
## 1 x o x o o x o o x o x x o x
## 2 o o x o x x o x x x x x x x
## 3 x x x o o o o o o o x x x o
## 4 x x x o o o o o o o x o o
## 5 x x x x x o o o o o o x o x
## 6 x x x x x o o o o o o o o x
## 7 x x x x x o o o o o o o o x
```

Try an AR(12) model and examine residuals

```
#Set up external regressors and dummy vars
library(forecast)
monthDummies<-seasonaldummy(logBeer)
externReg<-data.frame(t, t2, t3, t4, monthDummies)

ar12_poly<-arima(logBeer, order=c(12,0,0), xreg=externReg)
ar12_poly
```

```
##
## Call:
## arima(x = logBeer, order = c(12, 0, 0), xreg = externReg)
##
## Coefficients:
##          ar1          ar2          ar3          ar4          ar5          ar6          ar7          ar8
##      -0.0185  -0.0373   0.0592  -0.0600   0.0852   0.1086  -0.0873   0.0867
## s.e.    0.0460   0.0449   0.0449   0.0441   0.0443   0.0442   0.0443   0.0438
##          ar9          ar10         ar11         ar12  intercept           t           t2          t3          t4
##      0.1421  -0.0077   0.2301   0.2099         4.6792  -0.0009   0e+00    0    0
## s.e.   0.0440   0.0440   0.0440   0.0453         0.0781   0.0014   2e-04    0    0
##          Jan          Feb          Mar          Apr          May          Jun          Jul
##      -0.1923  -0.2566  -0.1788  -0.2787  -0.3181  -0.4344  -0.3506
## s.e.    0.0154   0.0170   0.0156   0.0165   0.0167   0.0156   0.0167
##          Aug          Sep          Oct          Nov
##      -0.3012  -0.2658  -0.1259  -0.0756
## s.e.    0.0165   0.0157   0.0171   0.0155
##
## sigma^2 estimated as 0.003593:  log likelihood = 629.73,  aic = -1203.47
```


We seem to be having trouble getting fits for the trend line, ask about this Monday, try just using the month dummies.

```
ar12<-arima(logBeer, order=c(12,0,0), xreg=monthDummies)
ar12
```

```
##
## Call:
## arima(x = logBeer, order = c(12, 0, 0), xreg = monthDummies)
##
## Coefficients:
##          ar1          ar2          ar3          ar4          ar5          ar6          ar7          ar8
##      0.0631   0.0283   0.1255  -0.0127   0.1286   0.1447  -0.0733   0.0983
## s.e.   0.0464   0.0458   0.0458   0.0457   0.0458   0.0460   0.0461   0.0459
```

```
##          ar9      ar10      ar11      ar12  intercept      Jan      Feb
##      0.1446 -0.0278  0.2102  0.1670      4.9720 -0.1925 -0.2574
## s.e.  0.0461  0.0461  0.0461  0.0471      0.2644  0.0152  0.0169
##          Mar      Apr      May      Jun      Jul      Aug      Sep
##     -0.1798 -0.2797 -0.3188 -0.4349 -0.3507 -0.3012 -0.2657
## s.e.  0.0154  0.0165  0.0167  0.0153  0.0167  0.0165  0.0155
##          Oct      Nov
##     -0.126 -0.0757
## s.e.  0.017  0.0153
##
## sigma^2 estimated as 0.003946:  log likelihood = 606.68,  aic = -1165.36
```


```
tsdiag(ar12, gof.lag=20)
```



```
#residuals
```


```
ar12_resid<-ts(residuals(ar12), frequency=12, start=c(1956,1))
plot(ar12_resid, main="AR 12 model Residuals from Logged Beer\nseasonal Means and 4th order poly fit at same t
points(y=ar12_resid, x=time(ar12_resid), pch=as.vector(season(ar12_resid)))
abline(h=0)
```

AR 12 model Residuals from Logged Beer seasonal Means and 4th order poly fit at same time




```
par(mfrow=c(2,1))
acf(ar12_resid)
pacf(ar12_resid)
```

Series ar12_resid



Series ar12_resid




```
pacf_acf<-data.frame(acfVal=acf(ar12_resid, plot=FALSE)$acf, pacfVal=pacf(ar12_resid, plot=FALSE)$acf)
#print(pacf_acf)
```

After we choose a model, run all of the diagnostic tests

```
chosenMod<-ar12
modelString<-"SARIMA(12,0,0)(0,0,0)[12]"


par(mfrow=c(1,1))
plot(residuals(chosenMod), main=paste("Residuals of Model", modelString))
abline(h=0)
```



Comment:

```
par(mfrow=c(1,1))
qqnorm(residuals(chosenMod), main=paste("Normal QQ Plot of Residuals from", modelString))
qqline(residuals(chosenMod))
```



Normal QQ Plot of Residuals from SARIMA(12,0,0)(0,0,0)[12]



Comment:

```
par(mfrow=c(1,1))  
acf(residuals(chosenMod), main=paste("ACF of Residuals from", modelString))
```

ACF of Residuals from SARIMA(12,0,0)(0,0,0)[12]



Comment:

```
shapiro.test(residuals(chosenMod))
```

##

```
## Shapiro-Wilk normality test
##
## data: residuals(chosenMod)
## W = 0.98657, p-value = 0.000348
```

Comment:

```
LB.test(chosenMod, lag=24)
```

```
##
## Box-Ljung test
##
## data: residuals from chosenMod
## X-squared = 30.485, df = 12, p-value = 0.002359
```

Comment:

Make the forecasts


Set up external regressor data frame

```
newMonthDummy<-seasonaldummy(beer_forecast)
```

Plot the model forecasts

```
library(TSA)
```

```
TSA::plot.Arima(chosenMod,n.ahead=24,n1=c(1956,1), newxreg=newMonthDummy,
type='b',ylab='Logged Beer Production (Megalitres)',xlab='Year', col="red", lty=2, cex=.75)
```




```

predictions<-predict(chosenMod, newxreg=newMonthDummy, n.ahead=24)
pred<-predictions$pred
uci<-pred+2*predictions$se
lci<-pred-2*predictions$se

ymin=min(c(as.vector(lci),logBeer))-1
ymax=max(c(as.vector(uci),logBeer))+1
plot(logBeer,ylim=c(ymin,ymax),main=modelString, ylab='Logged Beer Production (Megalitres)')
lines(pred,col=2)
lines(uci,col=3)
lines(lci,col=3)

```

SARIMA(12,0,0)(0,0,0)[12]



```
ymin=min(c(as.vector(lci),logBeer))-1
ymax=max(c(as.vector(uci),logBeer))+1
plot(logBeer,xlim=c(1993, 1996), ylim=c(4.5,5.5),main=modelString, ylab='Logged Beer Production (Megalitres)')
lines(pred,col=2)
lines(uci,col=3)
lines(lci,col=3)
lines(log(beer_forecast), col="black")
```

SARIMA(12,0,0)(0,0,0)[12]

